

# Data Mining

## Classification Trees (2)

Ad Feelders

Universiteit Utrecht

# Basic Tree Construction Algorithm (control flow)

## Construct tree

nodelist  $\leftarrow \{\{\text{training data}\}\}$

Repeat

    current node  $\leftarrow$  select node from nodelist

    nodelist  $\leftarrow$  nodelist  $-$  current node

    if  $\text{impurity}(\text{current node}) > 0$

    then

$S \leftarrow$  set of candidate splits in current node

$s^* \leftarrow \arg \max_{s \in S} \text{impurity reduction}(s, \text{current node})$

        child nodes  $\leftarrow \text{apply}(s^*, \text{current node})$

        nodelist  $\leftarrow$  nodelist  $\cup$  child nodes

    fi

Until nodelist =  $\emptyset$

# Overfitting and Pruning

- The tree growing algorithm continues splitting until all leaf nodes of  $T$  contain examples of a single class (i.e. resubstitution error  $R(T) = 0$ ).
- Is this a good tree for predicting the class of new examples?
- Not unless the problem is truly “deterministic”!
- Problem of *overfitting*.

# Proposed Solutions

How can we prevent overfitting?


- Stopping Rules: e.g. don't expand a node if the impurity reduction of the best split is below some threshold.
- Pruning: grow a very large tree  $T_{\max}$  and merge back nodes.

Note: in the practical assignment we do use a stopping rule based on the `nmin` and `minleaf` parameters.

# Stopping Rules

Disadvantage: sometimes you first have to make a weak split to be able to follow up with a good split.


Since we only look one step ahead we may miss the good follow-up split.



# Pruning


- To avoid the problem of stopping rules, we first grow a very large tree on the training sample, and then *prune* this large tree.
- Objective: select the pruned subtree that has lowest *true* error rate.
- Problem: how to find this pruned subtree?
- Cost-complexity pruning (Breiman et al.; CART), also called *weakest link* pruning.

## Terminology: Tree $T$




$\tilde{T}$  denotes the collection of leaf nodes of tree  $T$ .  
 $\tilde{T} = \{t_5, t_6, t_7, t_8, t_9\}, |\tilde{T}| = 5$

## Terminology: Pruning $T$ in node $t_2$






Terminology:  $T$  after pruning in  $t_2$ :  $T - T_{t_2}$



## Terminology: Branch $T_{t_2}$



$$\tilde{T}_{t_2} = \{t_5, t_8, t_9\}, |\tilde{T}_{t_2}| = 3$$

# Cost-complexity pruning

- The total number of pruned subtrees of a balanced binary tree with  $\ell$  leaves is

$$\lfloor 1.5028369^\ell \rfloor$$

- With just 40 leaf nodes we have approximately 12 million pruned subtrees.
- Exhaustive search not recommended.
- Basic idea of cost-complexity pruning: reduce the number of pruned subtrees we have to consider by selecting the ones that are the “best of their kind” (in a sense to be defined shortly...)

## Total cost of a tree

Strike a balance between fit and complexity. Total cost  $C_\alpha(T)$  of tree  $T$

$$C_\alpha(T) = R(T) + \alpha|\tilde{T}|$$

Total cost consists of two components:

- resubstitution error  $R(T)$ , and
- a penalty for the complexity of the tree  $\alpha|\tilde{T}|, (\alpha \geq 0)$ .

Note:  $R(T) = \frac{\text{number of wrong classifications made by } T}{\text{number of examples in the training sample}}$

# Tree with lowest total cost

- Depending on the value of  $\alpha$ , different pruned subtrees will have the lowest total cost.
- For  $\alpha = 0$  (no complexity penalty) the tree with smallest resubstitution error wins.
- For higher values of  $\alpha$ , a less complex tree that makes a few more errors might win.

As it turns out, we can find a nested sequence of pruned subtrees of  $T_{\max}$ , such that the trees in the sequence minimize total cost for consecutive intervals of  $\alpha$  values.

# Smallest minimizing subtree

For any value of  $\alpha$ , there exists a smallest minimizing subtree  $T(\alpha)$  of  $T_{\max}$  that satisfies the following conditions:

- 1  $T(\alpha)$  minimizes total cost for that value of  $\alpha$ :  
$$C_{\alpha}(T(\alpha)) = \min_{T \leq T_{\max}} C_{\alpha}(T)$$
- 2  $T(\alpha)$  is a pruned subtree of all trees that minimize total cost:  
if  $C_{\alpha}(T) = C_{\alpha}(T(\alpha))$  then  $T(\alpha) \leq T$ .

**Note:**  $T' \leq T$  means  $T'$  is a pruned subtree of  $T$ ,  
i.e. it can be obtained by pruning  $T$  in 0 or more nodes.

# Sequence of subtrees

Construct a *decreasing sequence* of pruned subtrees of  $T_{\max}$

$$T_{\max} > T_1 > T_2 > T_3 > \dots > \{t_1\}$$

(where  $t_1$  is the root node of the tree) such that  $T_k$  is the smallest minimizing subtree for  $\alpha \in [\alpha_k, \alpha_{k+1})$ .

**Note:** From a computational viewpoint, the important property is that  $T_{k+1}$  is a pruned subtree of  $T_k$ . No backtracking is required.

# Decomposition of total cost

Total cost has an additive decomposition over the leaf nodes of a tree:


$$C_{\alpha}(T) = \sum_{t \in \tilde{T}} (R(t) + \alpha)$$

$R(t)$  is the number of errors we make in node  $t$  if we predict the majority class, divided by the total number of observations in the training sample.




# Effect on cost of pruning in node $t$

Before pruning in  $t$



$$C_\alpha(T_t) = \sum_{t' \in \tilde{T}_t} (R(t') + \alpha)$$

After pruning in  $t$



$$C_\alpha(\{t\}) = R(t) + \alpha$$

## Finding the $T_k$ and corresponding $\alpha_k$

$T_t$ : branch of  $T$  with root node  $t$ .

After pruning in  $t$ , its contribution to total cost is:

$$C_\alpha(\{t\}) = R(t) + \alpha,$$

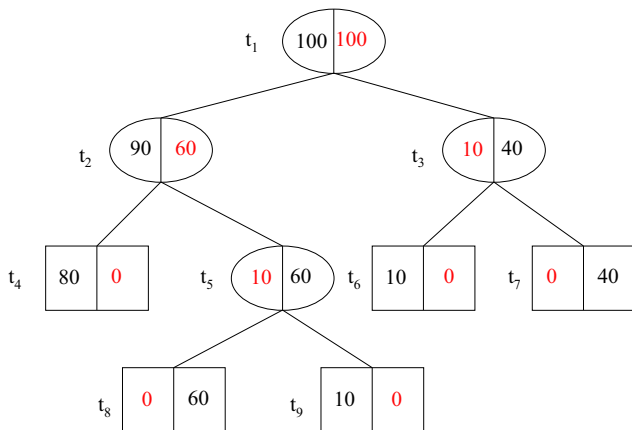
The contribution of  $T_t$  to the total cost is:

$$C_\alpha(T_t) = \sum_{t' \in \tilde{T}_t} (R(t') + \alpha) = R(T_t) + \alpha|\tilde{T}_t|$$

$T - T_t$  becomes better than  $T$  when

$$C_\alpha(\{t\}) = C_\alpha(T_t)$$

# Computing contributions to total cost of $T$



$$C_\alpha(\{t_2\}) = R(t_2) + \alpha = \frac{3}{10} + \alpha$$

$$C_\alpha(T_{t_2}) = R(T_{t_2}) + \alpha|\tilde{T}_{t_2}| = \alpha|\tilde{T}_{t_2}| + \sum_{t' \in \tilde{T}_{t_2}} R(t') = 3\alpha + 0$$

## Solving for $\alpha$

The total cost of  $T$  and  $T - T_t$  become equal when

$$C_\alpha(\{t\}) = C_\alpha(T_t),$$

At what value of  $\alpha$  does this happen?

$$R(t) + \alpha = R(T_t) + \alpha|\tilde{T}_t|$$

Solving for  $\alpha$  we get

$$\alpha = \frac{R(t) - R(T_t)}{|\tilde{T}_t| - 1}$$

Note: for this value of  $\alpha$  total cost of  $T$  and  $T - T_t$  is the same, but  $T - T_t$  is preferred because we want the *smallest* minimizing subtree.

## Computing $g(t)$ : the “critical” $\alpha$ value for node $t$

- For each non-terminal node  $t$  we compute its “critical” *alpha* value:

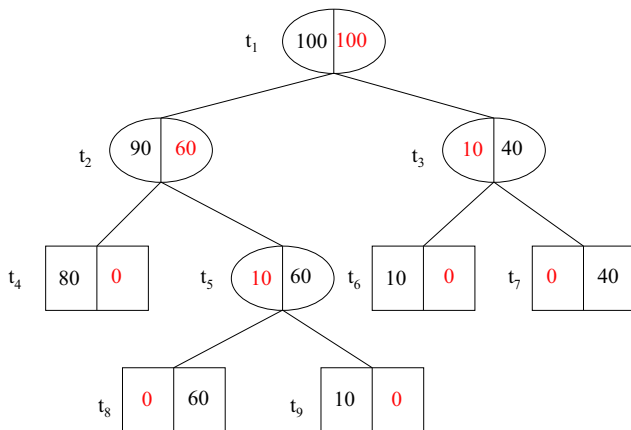
$$g(t) = \frac{R(t) - R(T_t)}{|\tilde{T}_t| - 1}$$

In words:

$$g(t) = \frac{\text{increase in error due to pruning in } t}{\text{decrease in } \# \text{ leaf nodes due to pruning in } t}$$

- Subsequently, we prune in the nodes for which  $g(t)$  is the smallest (the “weakest links”).
- This process is repeated until we reach the root node.

# Computing $g(t)$ : the “critical” $\alpha$ value for node $t$



$$g(t_1) = \frac{1}{8}, g(t_2) = \frac{3}{20}, g(t_3) = \frac{1}{20}, g(t_5) = \frac{1}{20}.$$

## Computing $g(t)$ : the “critical” $\alpha$ value for node $t$

Calculation examples:


$$g(t_1) = \frac{R(t_1) - R(T_{t_1})}{|\tilde{T}_{t_1}| - 1} = \frac{1/2 - 0}{5 - 1} = \frac{1}{8}$$

$$g(t_2) = \frac{R(t_2) - R(T_{t_2})}{|\tilde{T}_{t_2}| - 1} = \frac{3/10 - 0}{3 - 1} = \frac{3}{20}$$

$$g(t_3) = \frac{R(t_3) - R(T_{t_3})}{|\tilde{T}_{t_3}| - 1} = \frac{1/20 - 0}{2 - 1} = \frac{1}{20}$$

$$g(t_5) = \frac{R(t_5) - R(T_{t_5})}{|\tilde{T}_{t_5}| - 1} = \frac{1/20 - 0}{2 - 1} = \frac{1}{20}$$


# Finding the weakest links



$$g(t_1) = \frac{1}{8}, g(t_2) = \frac{3}{20}, g(t_3) = \frac{1}{20}, g(t_5) = \frac{1}{20}.$$




# Pruning in the weakest links



By pruning the weakest links we obtain the next tree in the sequence.

## Repeating the same procedure



$$g(t_1) = \frac{2}{10}, g(t_2) = \frac{1}{4}.$$

## Computing $g(t)$ : the “critical” $\alpha$ value for node $t$

Calculation examples:

$$g(t_1) = \frac{R(t_1) - R(T_{t_1})}{|\tilde{T}_{t_1}| - 1} = \frac{1/2 - 1/10}{3 - 1} = \frac{2}{10}$$

$$g(t_2) = \frac{R(t_2) - R(T_{t_2})}{|\tilde{T}_{t_2}| - 1} = \frac{3/10 - 1/20}{2 - 1} = \frac{1}{4}$$

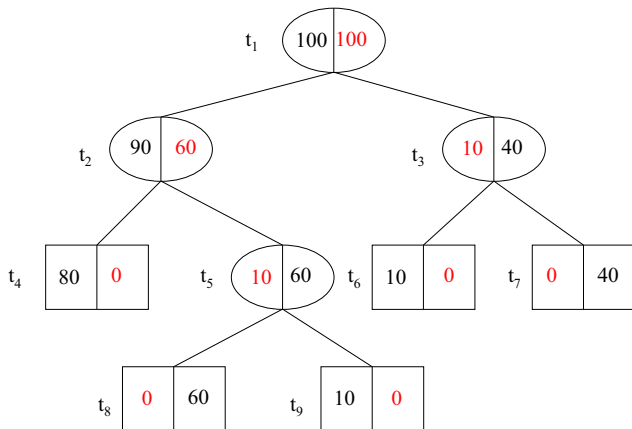
# Going back to the root

$t_1$

100	100
-----	-----


We have arrived at the root so we're done.

The best tree for  $\alpha \in [0, \frac{1}{20})$



The big tree is the best for values of  $\alpha$  below  $\frac{1}{20}$ .

The best tree for  $\alpha \in [\frac{1}{20}, \frac{2}{10})$



When  $\alpha$  reaches  $\frac{1}{20}$  this tree becomes the best.

The best tree for  $\alpha \in [\frac{2}{10}, \infty)$

$t_1$	<table><tr><td>100</td><td>100</td></tr></table>	100	100
100	100		

When  $\alpha$  reaches  $\frac{2}{10}$  the root wins and we're done.

# Computing the Pruning Sequence

```
 $T_1 \leftarrow T(\alpha = 0); \alpha_1 \leftarrow 0; k \leftarrow 1$   
While  $T_k > \{t_1\}$  do  
  For all non-terminal nodes  $t \in T_k$   
     $g_k(t) \leftarrow \frac{R(t) - R(T_{k,t})}{|\tilde{T}_{k,t}| - 1}$   
   $\alpha_{k+1} \leftarrow \min_t g_k(t)$   
  Visit the nodes in post-order and prune  
  whenever  $g_k(t) = \alpha_{k+1}$  to obtain  $T_{k+1}$   
   $k \leftarrow k + 1$   
od
```

Note:  $T_{k,t}$  is the branch of  $T_k$  with root node  $t$ ,  
and  $T_k$  is the pruned tree in iteration  $k$ .



## Algorithm to compute $T_1$ from $T_{\max}$

If we don't continue splitting until all nodes are pure, then  $T_1 = T(\alpha = 0)$  may not be the same as  $T_{\max}$ .

**Compute  $T_1$  from  $T_{\max}$**

$T' \leftarrow T_{\max}$

Repeat

Pick any pair of terminal nodes  $\ell$  and  $r$   
with common parent  $t$  in  $T'$

such that  $R(t) = R(\ell) + R(r)$ , and set

$T' \leftarrow T' - T_t$  (i.e. prune  $T'$  in  $t$ )

Until no more such pair exists

$T_1 \leftarrow T'$

## Selection of the final tree: using a test set

Pick the tree  $T$  from the sequence with the lowest error rate  $R^{ts}(T)$  on the test set.


This is an *estimate* of the true error rate  $R^*(T)$  of  $T$ .

The standard error of this estimate is

$$SE(R^{ts}) = \sqrt{\frac{R^{ts}(1 - R^{ts})}{n_{test}}},$$

where  $n_{test}$  is the number of observations in the test set.

# Selection of the final tree: the 1-SE rule



1-SE rule: select the smallest tree with  $R^{ts}$  within one standard error of the minimum.

# Cross-Validation

- When the data set is relatively small, it is a bit of a waste to set aside part of the data for testing.
- A way to avoid this problem is to use *cross-validation*.

# Cross-Validation

- 1 Divide data into  $v$  folds.
- 2 Train on  $v - 1$  folds.
- 3 Predict on the remaining fold.
- 4 Leave out each of the  $v$  folds in turn.

First iteration:

fold	$X$	$Y$	$\hat{Y}$
1			
2			
3			
4			
5			$\hat{Y}^{(5)}$

# Cross-Validation

Second iteration:

fold	$X$	$Y$	$\hat{Y}$
1			
2			
3			
4			$\hat{Y}^{(4)}$
5			

# Cross-Validation

Third iteration:

fold	$X$	$Y$	$\hat{Y}$
1			$\hat{Y}^{(3)}$
2			
3			
4			
5			

# Cross-Validation

Fourth iteration:

fold	$X$	$Y$	$\hat{Y}$
1			$\hat{Y}^{(2)}$
2			
3			
4			
5			



# Cross-Validation

Fifth iteration:

fold	$X$	$Y$	$\hat{Y}$
1			$\hat{Y}^{(1)}$
2			
3			
4			
5			

# Cross-Validation

In the end we have out-of-sample predictions for all cases!

fold	$X$	$Y$	$\hat{Y}$
1			$\hat{Y}^{(1)}$
2			$\hat{Y}^{(2)}$
3			$\hat{Y}^{(3)}$
4			$\hat{Y}^{(4)}$
5			$\hat{Y}^{(5)}$

- 1 Perform cross-validation for different hyper-parameter settings (e.g.  $n_{min}$  and  $minleaf$ ).
- 2 Compute prediction error for each parameter setting.
- 3 Pick setting with lowest error.
- 4 Train with selected setting on complete data set.

# $\nu$ -fold cross-validation (general)

Let  $C$  be a complexity parameter of a learning algorithm (like  $\alpha$  in the classification tree algorithm). To select the best value of  $C$  from a range of values  $c_1, \dots, c_m$  we proceed as follows.

- ➊ Divide the data into  $\nu$  groups  $G_1, \dots, G_\nu$ .
- ➋ For each value  $c_i$  of  $C$ 
  - ➊ For  $j = 1, \dots, \nu$ 
    - ➊ Train with  $C = c_i$  on all data *except* group  $G_j$ .
    - ➋ Predict on group  $G_j$ .
  - ➋ Compute the CV prediction error for  $C = c_i$ .
- ➌ Select the value  $c^*$  of  $C$  with the smallest CV prediction error.
- ➍ Train on the complete training sample with  $C = c^*$

# Using cross-validation: Step 1

Grow a tree on the full data set, and compute  $\alpha_1, \alpha_2, \dots, \alpha_K$  and  $T_1 > T_2 > \dots > T_K$ .

Recall that  $T_k$  is the smallest minimizing subtree for  $\alpha \in [\alpha_k, \alpha_{k+1})$ .

Estimate the error of a tree  $T_k$  from this sequence as follows.

Set

$$\beta_1 = 0,$$

$$\beta_2 = \sqrt{\alpha_2 \alpha_3},$$

$$\beta_3 = \sqrt{\alpha_3 \alpha_4},$$

$$\dots,$$

$$\beta_{K-1} = \sqrt{\alpha_{K-1} \alpha_K},$$

$$\beta_K = \infty.$$

$\beta_k$  is the “representative” value for  $[\alpha_k, \alpha_{k+1})$ .

## Using cross-validation: Step 2

Divide the data set into  $v$  groups  $G_1, G_2, \dots, G_v$  (of approximately equal size) and for each group  $G_j$

- 1 Grow a tree on all data *except*  $G_j$ , and determine the smallest minimizing subtrees  $T^{(j)}(\beta_1), T^{(j)}(\beta_2), \dots, T^{(j)}(\beta_K)$  for this reduced data set.
- 2 Compute the error of  $T^{(j)}(\beta_k)$  ( $k = 1, \dots, K$ ) on  $G_j$ .

## Using cross-validation: Step 3

- ① For each  $\beta_k$ , sum the errors of  $T^{(j)}(\beta_k)$  over  $G_j$  ( $j = 1, \dots, v$ ).
- ② Let  $\beta_h$  be the one with the lowest overall error.  
Select  $T_h$  as the best tree.
- ③ Use the error rate computed with cross-validation as an estimate of its error rate.

Remark: Alternatively, we could again use the 1-SE rule in the final step to select the final tree from the sequence.

# Using cross-validation: Step 1

Tree sequence constructed on *full* data set:

- $T_1$  is the best tree for  $\alpha \in [0, \frac{1}{20})$ .
- $T_2$  is the best tree for  $\alpha \in [\frac{1}{20}, \frac{2}{10})$ .
- $T_3$  is the best tree for  $\alpha \in [\frac{2}{10}, \infty)$ .

Set

$\beta_1 = 0$ , value corresponding to  $T_1$

$\beta_2 = \sqrt{\frac{1}{20} \frac{2}{10}} = \frac{1}{10}$ , value corresponding to  $T_2$

$\beta_3 = \infty$ , value corresponding to  $T_3$  (root).

## Using cross-validation: Step 2

Divide the data set in  $v = 4$  groups  $G_1, G_2, G_3, G_4$  of size 50 each.

First CV-run

- 1 Build a tree on all data *except*  $G_1$ , and determine the smallest minimizing subtrees  $T^{(1)}(0)$ ,  $T^{(1)}(\frac{1}{10})$  and  $T^{(1)}(\infty)$ .
- 2 Compute the error of those trees on  $G_1$ .

Repeat this procedure for  $G_2$ ,  $G_3$  and  $G_4$ .



## Using cross-validation: Step 3

CV-run	$\beta_1 = 0$	$\beta_2 = \frac{1}{10}$	$\beta_3 = \infty$
1	20	10	25
2	18	8	25
3	22	9	25
4	20	13	25
Total	80	40	100

$\beta_2$  wins (40 errors), so  $T_2$  gets selected.

We estimate the error rate of  $T_2$  at 20%.

# Building Trees in R: Rpart

Pima Indians Diabetes Database from the UCI ML Repository

1. Number of times pregnant
2. Plasma glucose concentration a 2 hours in an oral glucose tolerance test
3. Diastolic blood pressure (mm Hg)
4. Triceps skin fold thickness (mm)
5. 2-Hour serum insulin ( $\mu$ U/ml)
6. Body mass index ( $\text{weight in kg} / (\text{height in m})^2$ )
7. Diabetes pedigree function
8. Age (years)
9. Class variable (0 or 1)

Class Distribution: (class value 1 is interpreted as "tested positive for diabetes")

Class Value	Number of instances
0	500
1	268

# Building Trees in R: Rpart

```
> pima.dat[1:5,]  
  npreg plasma bp triceps serum  bmi pedigree age class  
1     6   148 72    35     0 33.6   0.627  50    1  
2     1    85 66    29     0 26.6   0.351  31    0  
3     8   183 64     0     0 23.3   0.672  32    1  
4     1    89 66    23    94 28.1   0.167  21    0  
5     0   137 40    35   168 43.1   2.288  33    1
```

```
> library(rpart)  
> pima.tree <- rpart(class ~.,data=pima.dat,cp=0,minbucket=1,minsplit=2,method="class")  
> printcp(pima.tree)
```

Classification tree:

```
rpart(formula = class ~ ., data = pima.dat, method = "class",  
      cp = 0, minbucket = 1, minsplit = 2)
```

Variables actually used in tree construction:

```
[1] age      bmi      bp      npreg    pedigree plasma  serum   triceps
```

Root node error: 268/768 = 0.34896

n= 768

# cptable: the pruning sequence


	CP	nsplit	rel error	xerror	xstd
1	0.2425373	0	1.000000	1.00000	0.049288
2	0.1044776	1	0.757463	0.80970	0.046558
3	0.0174129	2	0.652985	0.71269	0.044698
4	0.0149254	5	0.600746	0.74627	0.045381
5	0.0130597	9	0.541045	0.75000	0.045454
6	0.0111940	12	0.492537	0.72761	0.045007
7	0.0087065	16	0.447761	0.71642	0.044776
8	0.0074627	19	0.421642	0.72761	0.045007
9	0.0062189	23	0.391791	0.73134	0.045083
10	0.0055970	28	0.358209	0.73881	0.045233
11	0.0049751	42	0.272388	0.74254	0.045307
12	0.0044776	45	0.257463	0.74254	0.045307
13	0.0037313	50	0.235075	0.78731	0.046159
14	0.0027985	88	0.093284	0.79851	0.046360
15	0.0024876	96	0.070896	0.82463	0.046814
16	0.0018657	109	0.037313	0.82463	0.046814
17	0.0000000	128	0.000000	0.84701	0.047184

CP is  $\alpha$  divided by the resubstitution error in the root node.

Example: tree with 2 splits is best for  $CP \in [0.0174129, 0.1044776)$ .

Tree with 2 splits has cross-validation error of  $0.34896 \times 0.71269 = 0.2487$ .

# Plot of pruning sequence



# Selecting the best tree

```
> pima.pruned <- prune(pima.tree,cp=0.02)  
> post(pima.pruned)
```

